

Ridge Regression

Problem:

Generate a predictor vector \mathbf{X} of length $n = 100$ (random vector \mathbf{X}), as well as a noise vector $\boldsymbol{\epsilon}$ of length $n = 100$. Generate a response vector \mathbf{Y} of length $n = 100$ according to the following model:

$$Y_i = \beta_0 + \beta_1 X_i + \beta_2 X_i^2 + \beta_3 X_i^3 + \epsilon_i$$

where:

- $\beta_0 = 50$
- $\beta_1 = 10$
- $\beta_2 = -20$
- $\beta_3 = 0.1$

Perform ridge regression using \mathbf{X} , \mathbf{X}^2 , \mathbf{X}^3 , and \mathbf{X}^4 as predictors. Choose any two different values of λ (different from 0 and ∞). With each λ , perform ridge regression both **with** and **without** standardizing the predictors. Then, compare the results.

Note: No built-in functions are allowed.

```
In [1]: import numpy as np
```

```
In [2]: # set random seed for reproducibility
np.random.seed(403)
```

```
In [3]: # define parameters
n = 100
x = np.random.rand(n)
epsilon = np.random.normal(0, 1, n)
b0 = 50
b1 = 10
b2 = -20
b3 = 0.1
```

```
In [4]: # standardizing the data
x_mean = np.mean(x)
x_std = np.std(x)
x_standardized = (x - x_mean) / x_std
```

```
In [5]: # generate Y
Y = b0 + b1*x + b2*x**2 + b3*x**3 + epsilon
```

```
In [6]: def ridge_regression(X, Y, lmd):
        n, p = X.shape
        I = np.eye(p)
        beta_hat = np.linalg.inv(X.T @ X + lmd * I) @ X.T @ Y
        return beta_hat
```

```
In [7]: # not standardizing the data
        lmd0 = 0.1
        b_ridge0 = ridge_regression(np.column_stack((np.ones(n), x, x**2, x**3)), Y, lmd0)
        # predict using the estimated coefficients
        Y_hat0 = np.column_stack((np.ones(n), x, x**2, x**3)) @ b_ridge0
        # calculate the mse
        mse0 = np.mean((Y - Y_hat0)**2)
        print(f"mse with lambda = {lmd0} without standardizing the data: {mse0}")

        # standardizing the data
        b_ridge0_ = ridge_regression(np.column_stack((np.ones(n), x_standardized, x_standardized**2, x_standardized**3)), Y, lmd0)
        # predict using the estimated coefficients
        Y_hat0_ = np.column_stack((np.ones(n), x_standardized, x_standardized**2, x_standardized**3)) @ b_ridge0_
        # calculate the mse
        mse0_ = np.mean((Y - Y_hat0_)**2)
        print(f"mse with lambda = {lmd0} with standardizing the data: {mse0_}")
```

mse with lambda = 0.1 without standardizing the data: 1.0939062762083729

mse with lambda = 0.1 with standardizing the data: 1.0740449019290557

```
In [8]: lmd1 = 0.01
        # not standardizing the data
        b_ridge1 = ridge_regression(np.column_stack((np.ones(n), x, x**2, x**3)), Y, lmd1)
        # predict using the estimated coefficients
        Y_hat1 = np.column_stack((np.ones(n), x, x**2, x**3)) @ b_ridge1
        # calculate the mse
        mse1 = np.mean((Y - Y_hat1)**2)
        print(f"mse with lambda = {lmd1} without standardizing the data: {mse1}")

        # standardizing the data
        b_ridge1_ = ridge_regression(np.column_stack((np.ones(n), x_standardized, x_standardized**2, x_standardized**3)), Y, lmd1)
        # predict using the estimated coefficients
        Y_hat1_ = np.column_stack((np.ones(n), x_standardized, x_standardized**2, x_standardized**3)) @ b_ridge1_
        # calculate the mse
        mse1_ = np.mean((Y - Y_hat1_)**2)
        print(f"mse with lambda = {lmd1} with standardizing the data: {mse1_}")
```

mse with lambda = 0.01 without standardizing the data: 1.0742068107622686

mse with lambda = 0.01 with standardizing the data: 1.068984594996808

Conclusion

1. With both λ 0.1 and 0.01, the mse after standardizing X has dropped, indicating an improved performance of ridge regression with standardization.
2. By dropping λ from 0.1 to 0.01, the mse decreased, indicating a lower penalty on coefficients with less regularization, introducing smaller bias.

Lasso Regression

Problem:

Use the dataset you generated in Problem 1 and fit the model for the same set of predictors using **Lasso regression**.

Choose any two different values of λ (different from 0 and ∞).

With each λ , perform **Lasso regression** without standardizing the predictors. Then perform **Lasso regression** standardizing the predictors.

Questions:

- What can you conclude from these experiments?

Note: No built-in functions are allowed.

```
In [9]: x.shape, Y.shape, x_standardized.shape
```

```
Out[9]: ((100,), (100,), (100,))
```

```
In [10]: lmd0, lmd1
```

```
Out[10]: (0.1, 0.01)
```

```
In [11]: class lasso_reg:
    def __init__(self, lmd, tol=1e-4, max_iter=1000):
        self.lmd = lmd
        self.tol = tol
        self.max_iter = max_iter
        self.coef_ = None

    def fit(self, X, Y):
        n, p = X.shape
        b = np.zeros(p)
        b_old = np.zeros(p)

        X_t_X = X.T @ X

        for _ in range(self.max_iter):
            for j in range(p):
                # compute the partial residual
                residual = Y - X @ b + X[:, j] * b[j]

                # update coefficient using soft-thresholding
                rho = X[:, j].T @ residual
                b[j] = self._soft_threshold(rho / X_t_X[j, j], self.lmd)

            # early stop if converge
            if np.linalg.norm(b - b_old, ord=2) < self.tol:
                break

            b_old = b.copy()

        self.coef_ = b
```

```

def _soft_threshold(self, rho, lmd):
    if rho > lmd:
        return rho - lmd
    elif rho < -lmd:
        return rho + lmd
    else:
        return 0

def predict(self, X):
    return X @ self.coef_

def cal_mse(self, X, Y):
    Y_hat = self.predict(X)
    return np.mean((Y - Y_hat)**2)

```

```

In [12]: # case without standardizing the data
lasso0 = lasso_reg(lmd0)
lasso0.fit(np.column_stack((np.ones(n), x, x**2, x**3)), Y)
# calculate the mse
mse_lasso0 = lasso0.cal_mse(np.column_stack((np.ones(n), x, x**2, x**3)), Y)
print(f"mse with lambda = {lmd0} without standardizing the data: {mse_lasso0}")

lasso1 = lasso_reg(lmd1)
lasso1.fit(np.column_stack((np.ones(n), x, x**2, x**3)), Y)
# calculate the mse
mse_lasso1 = lasso1.cal_mse(np.column_stack((np.ones(n), x, x**2, x**3)), Y)
print(f"mse with lambda = {lmd1} without standardizing the data: {mse_lasso1}")

```

mse with lambda = 0.1 without standardizing the data: 1.1508305732256225
mse with lambda = 0.01 without standardizing the data: 1.0871402994665913

```

In [13]: # case with standardizing the data
lasso0_ = lasso_reg(lmd0)
lasso0_.fit(np.column_stack((np.ones(n), x_standardized, x_standardized**2, x_standardized**3)), Y)
# calculate the mse
mse_lasso0_ = lasso0_.cal_mse(np.column_stack((np.ones(n), x_standardized, x_standardized**2, x_standardized**3)), Y)
print(f"mse with lambda = {lmd0} with standardizing the data: {mse_lasso0_}")

lasso1_ = lasso_reg(lmd1)
lasso1_.fit(np.column_stack((np.ones(n), x_standardized, x_standardized**2, x_standardized**3)), Y)
# calculate the mse
mse_lasso1_ = lasso1_.cal_mse(np.column_stack((np.ones(n), x_standardized, x_standardized**2, x_standardized**3)), Y)
print(f"mse with lambda = {lmd1} with standardizing the data: {mse_lasso1_}")

```

mse with lambda = 0.1 with standardizing the data: 1.1755404017481759
mse with lambda = 0.01 with standardizing the data: 1.070824122036093

```

In [14]: # since unclear the effect of changing lambda, use train-test split to evaluate the model
x_train = x[:80]
x_test = x[80:]
y_train = Y[:80]
y_test = Y[80:]
# standardizing the data
x_train_standardized = (x_train - np.mean(x_train)) / np.std(x_train)
x_test_standardized = (x_test - np.mean(x_train)) / np.std(x_train)

```

```

In [15]: # case without standardizing the data
lasso0 = lasso_reg(lmd0)
lasso0.fit(np.column_stack((np.ones(80), x_train, x_train**2, x_train**3)), y_train)

```

```
# calculate train mse
mse_train_lasso0 = lasso0.cal_mse(np.column_stack((np.ones(80), x_train, x_train**2, x_train**3)), y_train)
# calculate test mse
mse_test_lasso0 = lasso0.cal_mse(np.column_stack((np.ones(20), x_test, x_test**2, x_test**3)), y_test)
print(f"train mse with lambda = {lmd0} without standardizing the data: {mse_train_lasso0}")
print(f"test mse with lambda = {lmd0} without standardizing the data: {mse_test_lasso0}")

lasso1 = lasso_reg(lmd1)
lasso1.fit(np.column_stack((np.ones(80), x_train, x_train**2, x_train**3)), y_train)
# calculate train mse
mse_train_lasso1 = lasso1.cal_mse(np.column_stack((np.ones(80), x_train, x_train**2, x_train**3)), y_train)
# calculate test mse
mse_test_lasso1 = lasso1.cal_mse(np.column_stack((np.ones(20), x_test, x_test**2, x_test**3)), y_test)
print(f"train mse with lambda = {lmd1} without standardizing the data: {mse_train_lasso1}")
print(f"test mse with lambda = {lmd1} without standardizing the data: {mse_test_lasso1}")
```

train mse with lambda = 0.1 without standardizing the data: 1.2612498747532779
test mse with lambda = 0.1 without standardizing the data: 0.6818476146548968
train mse with lambda = 0.01 without standardizing the data: 1.1930347305758442
test mse with lambda = 0.01 without standardizing the data: 0.6829818446647704

```
In [16]: # case with standardizing the data
lasso0_ = lasso_reg(lmd0)
lasso0_.fit(np.column_stack((np.ones(80), x_train_standardized, x_train_standardized**2, x_train_standardized**3)), y_train)
# calculate train mse
mse_train_lasso0_ = lasso0_.cal_mse(np.column_stack((np.ones(80), x_train_standardized, x_train_standardized**2, x_train_standardized**3)), y_train)
# calculate test mse
mse_test_lasso0_ = lasso0_.cal_mse(np.column_stack((np.ones(20), x_test_standardized, x_test_standardized**2, x_test_standardized**3)), y_test)
print(f"train mse with lambda = {lmd0} with standardizing the data: {mse_train_lasso0_}")
print(f"test mse with lambda = {lmd0} with standardizing the data: {mse_test_lasso0_}")

lasso1_ = lasso_reg(lmd1)
lasso1_.fit(np.column_stack((np.ones(80), x_train_standardized, x_train_standardized**2, x_train_standardized**3)), y_train)
# calculate train mse
mse_train_lasso1_ = lasso1_.cal_mse(np.column_stack((np.ones(80), x_train_standardized, x_train_standardized**2, x_train_standardized**3)), y_train)
# calculate test mse
mse_test_lasso1_ = lasso1_.cal_mse(np.column_stack((np.ones(20), x_test_standardized, x_test_standardized**2, x_test_standardized**3)), y_test)
print(f"train mse with lambda = {lmd1} with standardizing the data: {mse_train_lasso1_}")
print(f"test mse with lambda = {lmd1} with standardizing the data: {mse_test_lasso1_}")
```

train mse with lambda = 0.1 with standardizing the data: 1.2858833371163871
test mse with lambda = 0.1 with standardizing the data: 0.6007472024541338
train mse with lambda = 0.01 with standardizing the data: 1.1760703588835029
test mse with lambda = 0.01 with standardizing the data: 0.6446689662376079

Conclusion

1. With both λ 0.1 and 0.01, either test and train set, the mse after standardizing X has dropped, indicating an improved performance of lasso regression with standardization.
2. By dropping λ from 0.1 to 0.01, the mse deviate differently in train and test sets. MSE would decrease in train set and increase in test set when λ drops from 0.1 to 0.01. Indicating decrease λ would increase the overfitting problem.